



Mariia Mohylna

Non-portable kernel-based models

1

Key concepts
HIP vs CUDA
Hardware and grip
First program

2

Optimizations
Shared memory
Matrix multiplication
Parallel reduction

3

Honorable mentions
Memory types
Streams
...and more

HIP vs CUDA

- ✓ Focus **solely** and provide tools to **optimize** massive parallelization on GPUs

CUDA

- Supported on NVIDIA GPUs

HIP

- Supported on AMD and (theoretically) NVIDIA GPUs

HIP vs CUDA

- ✓ Focus **solely** and provide tools to **optimize** massive parallelization on GPUs

CUDA

- Supported on NVIDIA GPUs
- NVCC compiler

HIP

- Supported on AMD and (theoretically) NVIDIA GPUs
- HIPCC compiler

HIP vs CUDA

- ✓ Focus **solely** and provide tools to **optimize** massive parallelization on GPUs

CUDA

- Supported on NVIDIA GPUs
- NVCC compiler
- *.cuh or *.cu

HIP

- Supported on AMD and (theoretically) NVIDIA GPUs
- HIPCC compiler
- *.hip.hpp or *.hip.cpp

HIP vs CUDA

- ✓ Focus **solely** with and provide tools to **optimize** massive parallelization on GPUs

CUDA

- Supported on NVIDIA GPUs
- NVCC compiler
- *.cuh or *.cu

compulsory

HIP

- Supported on AMD and (theoretically) NVIDIA GPUs
- HIPCC compiler
- *.hip.hpp or *.hip.cpp

convinient

HIP vs CUDA

- ✓ Focus **solely** and provide tools to **optimize** massive parallelization on GPUs

CUDA

- Supported on NVIDIA GPUs
- NVCC compiler
- *.cuh or *.cu
- cudaMalloc, cudaMemcpy, cudaFree

HIP

- Supported on AMD and (theoretically) NVIDIA GPUs
- HIPCC compiler
- *.hip.hpp or *.hip.cpp
- hipMalloc, hipMemcpy, hipFree

HIP vs CUDA

- ✓ Focus **solely** and provide tools to **optimize** massive parallelization on GPUs

CUDA

- Supported on NVIDIA GPUs
- NVCC compiler
- *.cuh or *.cu
- cudaMalloc, cudaMemcpy, cudaFree
- #include<cuda_runtime.h>

HIP

- Supported on AMD and (theoretically) NVIDIA GPUs
- HIPCC compiler
- *.hip.hpp or *.hip.cpp
- hipMalloc, hipMemcpy, hipFree
- #include<hip/hip_runtime.h>

HIP vs CUDA

- ✓ Focus **solely** and provide tools to **optimize** massive parallelization on GPUs

CUDA


- Supported on NVIDIA GPUs
- NVCC compiler
- *.cuh or *.cu
- cudaMalloc, cudaMemcpy, cudaFree
- #include<cuda_runtime.h>

HIP

- Supported on AMD and (theoretically) NVIDIA GPUs
- HIPCC compiler
- *.hip.hpp or *.hip.cpp
- hipMalloc, hipMemcpy, hipFree
- #include<hip/hip_runtime.h>

✓ Mostly have the same functionality, but better to make sure

Interlude: rank addressing





Key concepts

Optimizations

Honorable mentions

Interlude: rank addressing




Key concepts


Optimizations

Honorable mentions


Interlude: rank addressing





Interlude: rank addressing



Interlude: rank addressing




Interlude: rank addressing



$idx = iy * nx + ix$

“global” rank




Key concepts

Optimizations

Honorable mentions

Interlude: rank addressing





Key concepts

Optimizations

Honorable mentions

Interlude: rank addressing




$$idx = iy * nx + ix$$

➤ Reverse:


$$iy = idx / nx$$
$$ix = idx \% nx$$

Hardware and grid




- ✓ Do not have program counters
- ✓ FP32/int

Hardware and grid




Hardware and grid




- ✓ Share a single program counter
- ✓ Add FP64

Hardware and grid




Hardware and grid



Hardware and grid

HIP




- ✓ Thousands cores
- ✓ Global memory

Key concepts

Optimizations

Honorable mentions

Hardware and grid




Key concepts


Optimizations

Honorable mentions


Hardware and grid



Hardware and grid




Hardware and grid




... and more about grid

- ✓ Thread rank in block: **threadIdx.x**
- ✓ Block rank in grid: **blockIdx.x**
- ✓ Number of threads in block: **blockDim.x**
- ✓ Number of blocks in grid: **gridDim.x**

"global" thread rank




$$\text{tid} = \text{blockIdx.x} * \text{blockDim.x} + \text{threadIdx.x}$$



... and more about grid


- ✓ Thread rank in block: **threadIdx.x**
- ✓ Block rank in grid: **blockIdx.x**
- ✓ Number of threads in block: **blockDim.x**
- ✓ Number of blocks in grid: **gridDim.x**



"global" thread rank


$$\text{tid} = \text{blockIdx.x} * \text{blockDim.x} + \text{threadIdx.x}$$

- Max number of threads in block: 1024
- Must be a multiple of 32 (64)



... and more about grid


- ✓ Thread rank in block: **threadIdx.x**
- ✓ Block rank in grid: **blockIdx.x**
- ✓ Number of threads in block: **blockDim.x**
- ✓ Number of blocks in grid: **gridDim.x**



"global" thread rank


$$\text{tid} = \text{blockIdx.x} * \text{blockDim.x} + \text{threadIdx.x}$$

- Max number of threads in block: 1024
- Must be a multiple of 32 (64)
- Threads and blocks can be packed 2D ...



... and more about grid


- ✓ Thread rank in block: **threadIdx.x**
- ✓ Block rank in grid: **blockIdx.x**
- ✓ Number of threads in block: **blockDim.x**
- ✓ Number of blocks in grid: **gridDim.x**



"global" thread rank


$$\text{tid} = \text{blockIdx.x} * \text{blockDim.x} + \text{threadIdx.x}$$

- Max number of threads in block: 1024
- Must be a multiple of 32 (64)
- Threads and blocks can be packed 2D and 3D



... and more about grid

- ✓ Thread rank in block: **threadIdx.x**
- ✓ Block rank in grid: **blockIdx.x**
- ✓ Number of threads in block: **blockDim.x**
- ✓ Number of blocks in grid: **gridDim.x**



"global" thread rank

$$\text{tid} = \text{blockIdx.x} * \text{blockDim.x} + \text{threadIdx.x}$$

- Max number of threads in block: 1024
- Must be a multiple of 32 (64)



- Threads and blocks can be packed 2D and 3D

2D threads on 2D blocks:

$$\begin{aligned}\text{blockId} &= (\text{gridDim.x} * \text{blockIdx.y}) + \text{blockIdx.x} \\ \text{tid} &= (\text{blockId} * (\text{blockDim.x} * \text{blockDim.y})) \\ &\quad + (\text{threadIdx.y} * \text{blockDim.x}) \\ &\quad + \text{threadIdx.x}\end{aligned}$$

... and more about grid

- ✓ Thread rank in block: **threadIdx.x**
- ✓ Block rank in grid: **blockIdx.x**
- ✓ Number of threads in block: **blockDim.x**
- ✓ Number of blocks in grid: **gridDim.x**



"global" thread rank

$$tid = blockIdx.x * blockDim.x + threadIdx.x$$

Fastest running index –
neighbouring threads in memory

$$\begin{aligned} blockId &= (gridDim.x * blockIdx.y) + blockIdx.x \\ tid &= (blockId * (blockDim.x * blockDim.y)) \\ &\quad + (threadIdx.y * blockDim.x) \\ &\quad + threadIdx.x \end{aligned}$$

Data transfer

```
const int N = 10000;
float *Ah, *Bh, *Ch, *Cref;
float *Ad, *Bd, *Cd;

Ah = (float *)malloc(N * sizeof(float)); Bh = (float *)malloc(N * sizeof(float));
Ch = (float *)malloc(N * sizeof(float)); Cref = (float *)malloc(N * sizeof(float));

for (int i = 0; i < N; i++) {
    Ah[i] = sin(i) * 2.3;
    Bh[i] = cos(i) * 1.1;
    Cref[i] = Ah[i] + Bh[i];
}

hipMalloc((void **)&Ad, N * sizeof(float));
hipMalloc((void **)&Bd, N * sizeof(float));
hipMalloc((void **)&Cd, N * sizeof(float));

hipMemcpy(Ad, Ah, sizeof(float) * N, hipMemcpyHostToDevice);
hipMemcpy(Bd, Bh, sizeof(float) * N, hipMemcpyHostToDevice);

dim3 blocks, threads;
threads = dim3(256, 1, 1);
blocks = dim3((N + 256 - 1) / 256, 1, 1);

// hipLaunchKernelGGI(vector_add, blocks, threads, 0, 0, Ad, Bd, Cd, N); // or
vector_add<<<blocks, threads, 0>>>(Ad, Bd, Cd, N);

hipMemcpy(Ch, Cd, sizeof(float) * N, hipMemcpyDeviceToHost);

hipFree(Ad); hipFree(Bd); hipFree(Cd);
free(Ah); free(Bh); free(Ch); free(Cref);
```

- CPU and GPU **cannot access** and modify each other's memory directly
- But they **can copy** data from each other

Data transfer

```
const int N = 10000;
float *Ah, *Bh, *Ch, *Cref;
float *Ad, *Bd, *Cd;

Ah = (float *)malloc(N * sizeof(float)); Bh = (float *)malloc(N * sizeof(float));
Ch = (float *)malloc(N * sizeof(float)); Cref = (float *)malloc(N * sizeof(float));

for (int i = 0; i < N; i++) {
    Ah[i] = sin(i) * 2.3;
    Bh[i] = cos(i) * 1.1;
    Cref[i] = Ah[i] + Bh[i];
}

hipMalloc((void **)&Ad, N * sizeof(float));
hipMalloc((void **)&Bd, N * sizeof(float));
hipMalloc((void **)&Cd, N * sizeof(float));

hipMemcpy(Ad, Ah, sizeof(float) * N, hipMemcpyHostToDevice);
hipMemcpy(Bd, Bh, sizeof(float) * N, hipMemcpyHostToDevice);

dim3 blocks, threads;
threads = dim3(256, 1, 1);
blocks = dim3((N + 256 - 1) / 256, 1, 1);

// hipLaunchKernelGGL(vector_add, blocks, threads, 0, 0, Ad, Bd, Cd, N); // or
vector_add<<<blocks, threads, 0>>>(Ad, Bd, Cd, N);

hipMemcpy(Ch, Cd, sizeof(float) * N, hipMemcpyDeviceToHost);

hipFree(Ad); hipFree(Bd); hipFree(Cd);
free(Ah); free(Bh); free(Ch); free(Cref);
```

- CPU and GPU **cannot access** and modify each other's memory directly
- But they **can copy** data from each other

h = host = CPU
d = device = GPU

Data transfer

```

const int N = 10000;
float *Ah, *Bh, *Ch, *Cref;
float *Ad, *Bd, *Cd;

Ah = (float *)malloc(N * sizeof(float)); Bh = (float *)malloc(N * sizeof(float));
Ch = (float *)malloc(N * sizeof(float)); Cref = (float *)malloc(N * sizeof(float));

for (int i = 0; i < N; i++) {
    Ah[i] = sin(i) * 2.3;
    Bh[i] = cos(i) * 1.1;
    Cref[i] = Ah[i] + Bh[i];
}

hipMalloc((void **)&Ad, N * sizeof(float));
hipMalloc((void **)&Bd, N * sizeof(float));
hipMalloc((void **)&Cd, N * sizeof(float));

hipMemcpy(Ad, Ah, sizeof(float) * N, hipMemcpyHostToDevice);
hipMemcpy(Bd, Bh, sizeof(float) * N, hipMemcpyHostToDevice);

dim3 blocks, threads;
threads = dim3(256, 1, 1);
blocks = dim3((N + 256 - 1) / 256, 1, 1);


// hipLaunchKernelGGL(vector_add, blocks, threads, 0, 0, Ad, Bd, Cd, N); // or
vector_add<<<blocks, threads, 0>>>(Ad, Bd, Cd, N);

hipMemcpy(Ch, Cd, sizeof(float) * N, hipMemcpyDeviceToHost);

hipFree(Ad); hipFree(Bd); hipFree(Cd);
free(Ah); free(Bh); free(Ch); free(Cref);

```

Allocate CPU and GPU memory



- CPU and GPU **cannot access** and modify each other's memory directly
- But they **can copy** data from each other

h = host = CPU
d = device = GPU

Data transfer

```

const int N = 10000;
float *Ah, *Bh, *Ch, *Cref;
float *Ad, *Bd, *Cd;

Ah = (float *)malloc(N * sizeof(float)); Bh = (float *)malloc(N * sizeof(float));
Ch = (float *)malloc(N * sizeof(float)); Cref = (float *)malloc(N * sizeof(float));

for (int i = 0; i < N; i++) {
    Ah[i] = sin(i) * 2.3;
    Bh[i] = cos(i) * 1.1;
    Cref[i] = Ah[i] + Bh[i];
}

hipMalloc((void **)&Ad, N * sizeof(float));
hipMalloc((void **)&Bd, N * sizeof(float));
hipMalloc((void **)&Cd, N * sizeof(float));

hipMemcpy(Ad, Ah, sizeof(float) * N, hipMemcpyHostToDevice);
hipMemcpy(Bd, Bh, sizeof(float) * N, hipMemcpyHostToDevice);

dim3 blocks, threads;
threads = dim3(256, 1, 1);
blocks = dim3((N + 256 - 1) / 256, 1, 1);

// hipLaunchKernelGGL(vector_add, blocks, threads, 0, 0, Ad, Bd, Cd, N); // or
vector_add<<<blocks, threads, 0, 0>>>(Ad, Bd, Cd, N);

hipMemcpy(Ch, Cd, sizeof(float) * N, hipMemcpyDeviceToHost);

hipFree(Ad); hipFree(Bd); hipFree(Cd);
free(Ah); free(Bh); free(Ch); free(Cref);

```

Allocate CPU and GPU memory

Copy data to and from GPU

- CPU and GPU **cannot access** and modify each other's memory directly
- But they **can copy** data from each other

h = host = CPU
d = device = GPU

- ✓ *hipMemcpy*: where, from where, how much data, H2D or D2H
- ✓ It is a **blocking** instruction

Data transfer

```

const int N = 10000;
float *Ah, *Bh, *Ch, *Cref;
float *Ad, *Bd, *Cd;

Ah = (float *)malloc(N * sizeof(float)); Bh = (float *)malloc(N * sizeof(float));
Ch = (float *)malloc(N * sizeof(float)); Cref = (float *)malloc(N * sizeof(float));

for (int i = 0; i < N; i++) {
    Ah[i] = sin(i) * 2.3;
    Bh[i] = cos(i) * 1.1;
    Cref[i] = Ah[i] + Bh[i];
}

hipMalloc((void **)&Ad, N * sizeof(float));
hipMalloc((void **)&Bd, N * sizeof(float));
hipMalloc((void **)&Cd, N * sizeof(float));

hipMemcpy(Ad, Ah, sizeof(float) * N, hipMemcpyHostToDevice);
hipMemcpy(Bd, Bh, sizeof(float) * N, hipMemcpyHostToDevice);

dim3 blocks, threads;
threads = dim3(256, 1, 1);
blocks = dim3((N + 256 - 1) / 256, 1, 1);

// hipLaunchKernelGGL(vector_add, blocks, threads, 0, 0, Ad, Bd, Cd, N); // or
vector_add<<<blocks, threads, 0, 0>>>(Ad, Bd, Cd, N);
hipDeviceSynchronize();
hipMemcpy(Ch, Cd, sizeof(float) * N, hipMemcpyDeviceToHost);

hipFree(Ad); hipFree(Bd); hipFree(Cd);
free(Ah); free(Bh); free(Ch); free(Cref);

```

Allocate CPU and GPU memory

Copy data to and from GPU

- CPU and GPU **cannot access** and modify each other's memory directly
- But they **can copy** data from each other

h = host = CPU
d = device = GPU

- ✓ *hipMemcpy*: where, from where, how much data, H2D or D2H
- ✓ It is a **blocking** instruction

Other blocking instruction:
hipDeviceSynchronize

Data transfer

```

const int N = 10000;
float *Ah, *Bh, *Ch, *Cref;
float *Ad, *Bd, *Cd;

Ah = (float *)malloc(N * sizeof(float)); Bh = (float *)malloc(N * sizeof(float));
Ch = (float *)malloc(N * sizeof(float)); Cref = (float *)malloc(N * sizeof(float));

for (int i = 0; i < N; i++) {
    Ah[i] = sin(i) * 2.3;
    Bh[i] = cos(i) * 1.1;
    Cref[i] = Ah[i] + Bh[i];
}

hipMalloc((void **)&Ad, N * sizeof(float));
hipMalloc((void **)&Bd, N * sizeof(float));
hipMalloc((void **)&Cd, N * sizeof(float));

hipMemcpy(Ad, Ah, sizeof(float) * N, hipMemcpyHostToDevice);
hipMemcpy(Bd, Bh, sizeof(float) * N, hipMemcpyHostToDevice);

dim3 blocks, threads;
threads = dim3(256, 1, 1);
blocks = dim3((N + 256 - 1) / 256, 1, 1);

// hipLaunchKernelGGL(vector_add, blocks, threads, 0, 0, Ad, Bd, Cd, N); // or
vector_add<<<blocks, threads, 0, 0>>>(Ad, Bd, Cd, N);

hipMemcpy(Ch, Cd, sizeof(float) * N, hipMemcpyDeviceToHost);

hipFree(Ad); hipFree(Bd); hipFree(Cd);
free(Ah); free(Bh); free(Ch); free(Cref);

```

Allocate CPU and GPU memory

Copy data to and from GPU

Free memory when finished

- CPU and GPU **cannot access** and modify each other's memory directly
- But they **can copy** data from each other

h = host = CPU
d = device = GPU

- ✓ *hipMemcpy*: where, from where, how much data, H2D or D2H
- ✓ It is a **blocking** instruction

Kernel launching

```
const int N = 10000;
float *Ah, *Bh, *Ch, *Cref;
float *Ad, *Bd, *Cd;

Ah = (float *)malloc(N * sizeof(float)); Bh = (float *)malloc(N * sizeof(float));
Ch = (float *)malloc(N * sizeof(float)); Cref = (float *)malloc(N * sizeof(float));

for (int i = 0; i < N; i++) {
    Ah[i] = sin(i) * 2.3;
    Bh[i] = cos(i) * 1.1;
    Cref[i] = Ah[i] + Bh[i];
}

hipMalloc((void **)&Ad, N * sizeof(float));
hipMalloc((void **)&Bd, N * sizeof(float));
hipMalloc((void **)&Cd, N * sizeof(float));

hipMemcpy(Ad, Ah, sizeof(float) * N, hipMemcpyHostToDevice);
hipMemcpy(Bd, Bh, sizeof(float) * N, hipMemcpyHostToDevice);

dim3 blocks, threads;
threads = dim3(256, 1, 1);
blocks = dim3((N + 256 - 1) / 256, 1, 1);

// hipLaunchKernelGGL(vector_add, blocks, threads, 0, 0, Ad, Bd, Cd, N); // or
vector_add<<<blocks, threads, 0, 0>>>(Ad, Bd, Cd, N);

hipMemcpy(Ch, Cd, sizeof(float) * N, hipMemcpyDeviceToHost);

hipFree(Ad); hipFree(Bd); hipFree(Cd);
free(Ah); free(Bh); free(Ch); free(Cref);
```

Kernel launching

➤ Step 1: create a grid for your task

```
const int N = 10000;
float *Ah, *Bh, *Ch, *Cref;
float *Ad, *Bd, *Cd;

Ah = (float *)malloc(N * sizeof(float)); Bh = (float *)malloc(N * sizeof(float));
Ch = (float *)malloc(N * sizeof(float)); Cref = (float *)malloc(N * sizeof(float));

for (int i = 0; i < N; i++) {
    Ah[i] = sin(i) * 2.3;
    Bh[i] = cos(i) * 1.1;
    Cref[i] = Ah[i] + Bh[i];
}

hipMalloc((void **)&Ad, N * sizeof(float));
hipMalloc((void **)&Bd, N * sizeof(float));
hipMalloc((void **)&Cd, N * sizeof(float));

hipMemcpy(Ad, Ah, sizeof(float) * N, hipMemcpyHostToDevice);
hipMemcpy(Bd, Bh, sizeof(float) * N, hipMemcpyHostToDevice);

dim3 blocks, threads;
threads = dim3(256, 1, 1);
blocks = dim3((N + 256 - 1) / 256, 1, 1);

// hipLaunchKernelGGI(vector_add, blocks, threads, 0, 0, Ad, Bd, Cd, N); // or
vector_add<<<blocks, threads, 0, 0>>>(Ad, Bd, Cd, N);

hipMemcpy(Ch, Cd, sizeof(float) * N, hipMemcpyDeviceToHost);

hipFree(Ad); hipFree(Bd); hipFree(Cd);
free(Ah); free(Bh); free(Ch); free(Cref);
```

Kernel launching

```

const int N = 10000;
float *Ah, *Bh, *Ch, *Cref;
float *Ad, *Bd, *Cd;

Ah = (float *)malloc(N * sizeof(float)); Bh = (float *)malloc(N * sizeof(float));
Ch = (float *)malloc(N * sizeof(float)); Cref = (float *)malloc(N * sizeof(float));

for (int i = 0; i < N; i++) {
    Ah[i] = sin(i) * 2.3;
    Bh[i] = cos(i) * 1.1;
    Cref[i] = Ah[i] + Bh[i];
}

hipMalloc((void **)&Ad, N * sizeof(float));
hipMalloc((void **)&Bd, N * sizeof(float));
hipMalloc((void **)&Cd, N * sizeof(float));

hipMemcpy(Ad, Ah, sizeof(float) * N, hipMemcpyHostToDevice);
hipMemcpy(Bd, Bh, sizeof(float) * N, hipMemcpyHostToDevice);

dim3 blocks, threads;
threads = dim3(256, 1, 1);
blocks = dim3((N + 256 - 1) / 256, 1, 1);

// hipLaunchKernelGGL(vector_add, blocks, threads, 0, 0, Ad, Bd, Cd, N); // or
vector_add<<<blocks, threads, 0, 0>>>(Ad, Bd, Cd, N);

hipMemcpy(Ch, Cd, sizeof(float) * N, hipMemcpyDeviceToHost);

hipFree(Ad); hipFree(Bd); hipFree(Cd);
free(Ah); free(Bh); free(Ch); free(Cref);

```

➤ Step 1: create a grid for your task



Kernel launching

```

const int N = 10000;
float *Ah, *Bh, *Ch, *Cref;
float *Ad, *Bd, *Cd;

Ah = (float *)malloc(N * sizeof(float)); Bh = (float *)malloc(N * sizeof(float));
Ch = (float *)malloc(N * sizeof(float)); Cref = (float *)malloc(N * sizeof(float));

for (int i = 0; i < N; i++) {
    Ah[i] = sin(i) * 2.3;
    Bh[i] = cos(i) * 1.1;
    Cref[i] = Ah[i] + Bh[i];
}

hipMalloc((void **)&Ad, N * sizeof(float));
hipMalloc((void **)&Bd, N * sizeof(float));
hipMalloc((void **)&Cd, N * sizeof(float));

hipMemcpy(Ad, Ah, sizeof(float) * N, hipMemcpyHostToDevice);
hipMemcpy(Bd, Bh, sizeof(float) * N, hipMemcpyHostToDevice);

dim3 blocks, threads;
threads = dim3(256, 1, 1);
blocks = dim3((N + 256 - 1) / 256, 1, 1);


// hipLaunchKernelGGI(vector_add, blocks, threads, 0, 0, Ad, Bd, Cd, N); // or
vector_add<<<blocks, threads, 0>>>(Ad, Bd, Cd, N);

hipMemcpy(Ch, Cd, sizeof(float) * N, hipMemcpyDeviceToHost);

hipFree(Ad); hipFree(Bd); hipFree(Cd);
free(Ah); free(Bh); free(Ch); free(Cref);

```

➤ Step 1: create a grid for your task



Kernel launching

```

const int N = 10000;
float *Ah, *Bh, *Ch, *Cref;
float *Ad, *Bd, *Cd;

Ah = (float *)malloc(N * sizeof(float)); Bh = (float *)malloc(N * sizeof(float));
Ch = (float *)malloc(N * sizeof(float)); Cref = (float *)malloc(N * sizeof(float));

for (int i = 0; i < N; i++) {
    Ah[i] = sin(i) * 2.3;
    Bh[i] = cos(i) * 1.1;
    Cref[i] = Ah[i] + Bh[i];
}

hipMalloc((void **)&Ad, N * sizeof(float));
hipMalloc((void **)&Bd, N * sizeof(float));
hipMalloc((void **)&Cd, N * sizeof(float));

hipMemcpy(Ad, Ah, sizeof(float) * N, hipMemcpyHostToDevice);
hipMemcpy(Bd, Bh, sizeof(float) * N, hipMemcpyHostToDevice);

dim3 blocks, threads;
threads = dim3(256, 1, 1);
blocks = dim3((N + 256 - 1) / 256, 1, 1);


// hipLaunchKernelGGL(vector_add, blocks, threads, 0, 0, Ad, Bd, Cd, N); // or
vector_add<<<blocks, threads, 0, 0>>>(Ad, Bd, Cd, N);

hipMemcpy(Ch, Cd, sizeof(float) * N, hipMemcpyDeviceToHost);

hipFree(Ad); hipFree(Bd); hipFree(Cd);
free(Ah); free(Bh); free(Ch); free(Cref);

```

➤ Step 1: create a grid for your task



Kernel launching

```

const int N = 10000;
float *Ah, *Bh, *Ch, *Cref;
float *Ad, *Bd, *Cd;

Ah = (float *)malloc(N * sizeof(float)); Bh = (float *)malloc(N * sizeof(float));
Ch = (float *)malloc(N * sizeof(float)); Cref = (float *)malloc(N * sizeof(float));

for (int i = 0; i < N; i++) {
    Ah[i] = sin(i) * 2.3;
    Bh[i] = cos(i) * 1.1;
    Cref[i] = Ah[i] + Bh[i];
}

hipMalloc((void **)&Ad, N * sizeof(float));
hipMalloc((void **)&Bd, N * sizeof(float));
hipMalloc((void **)&Cd, N * sizeof(float));

hipMemcpy(Ad, Ah, sizeof(float) * N, hipMemcpyHostToDevice);
hipMemcpy(Bd, Bh, sizeof(float) * N, hipMemcpyHostToDevice);

dim3 blocks, threads;
threads = dim3(256, 1, 1);
blocks = dim3((N + 256 - 1) / 256, 1, 1);


// hipLaunchKernelGGL(vector_add, blocks, threads, 0, 0, Ad, Bd, Cd, N); // or
vector_add<<<blocks, threads, 0, 0>>>(Ad, Bd, Cd, N);

hipMemcpy(Ch, Cd, sizeof(float) * N, hipMemcpyDeviceToHost);

hipFree(Ad); hipFree(Bd); hipFree(Cd);
free(Ah); free(Bh); free(Ch); free(Cref);

```

➤ Step 1: create a grid for your task



Kernel launching

```

const int N = 10000;
float *Ah, *Bh, *Ch, *Cref;
float *Ad, *Bd, *Cd;

Ah = (float *)malloc(N * sizeof(float)); Bh = (float *)malloc(N * sizeof(float));
Ch = (float *)malloc(N * sizeof(float)); Cref = (float *)malloc(N * sizeof(float));

for (int i = 0; i < N; i++) {
    Ah[i] = sin(i) * 2.3;
    Bh[i] = cos(i) * 1.1;
    Cref[i] = Ah[i] + Bh[i];
}

hipMalloc((void **)&Ad, N * sizeof(float));
hipMalloc((void **)&Bd, N * sizeof(float));
hipMalloc((void **)&Cd, N * sizeof(float));

hipMemcpy(Ad, Ah, sizeof(float) * N, hipMemcpyHostToDevice);
hipMemcpy(Bd, Bh, sizeof(float) * N, hipMemcpyHostToDevice);

dim3 blocks, threads;
threads = dim3(256, 1, 1);
blocks = dim3((N + 256 - 1) / 256, 1, 1);


// hipLaunchKernelGGL(vector_add, blocks, threads, 0, 0, Ad, Bd, Cd, N); // or
vector_add<<<blocks, threads, 0, 0>>>(Ad, Bd, Cd, N);

hipMemcpy(Ch, Cd, sizeof(float) * N, hipMemcpyDeviceToHost);

hipFree(Ad); hipFree(Bd); hipFree(Cd);
free(Ah); free(Bh); free(Ch); free(Cref);

```

➤ Step 1: create a grid for your task



- ✓ Total $num_threads$ = cover whole data
- ✓ Select block size
- ✓ Number of blocks = $num_threads / block_size$
(round up)

Kernel launching

```

const int N = 10000;
float *Ah, *Bh, *Ch, *Cref;
float *Ad, *Bd, *Cd;

Ah = (float *)malloc(N * sizeof(float)); Bh = (float *)malloc(N * sizeof(float));
Ch = (float *)malloc(N * sizeof(float)); Cref = (float *)malloc(N * sizeof(float));

for (int i = 0; i < N; i++) {
    Ah[i] = sin(i) * 2.3;
    Bh[i] = cos(i) * 1.1;
    Cref[i] = Ah[i] + Bh[i];
}

hipMalloc((void **)&Ad, N * sizeof(float));
hipMalloc((void **)&Bd, N * sizeof(float));
hipMalloc((void **)&Cd, N * sizeof(float));

hipMemcpy(Ad, Ah, sizeof(float) * N, hipMemcpyHostToDevice);
hipMemcpy(Bd, Bh, sizeof(float) * N, hipMemcpyHostToDevice);

dim3 blocks, threads;
threads = dim3(256, 1, 1);
blocks = dim3((N + 256 - 1) / 256, 1, 1);


// hipLaunchKernelGGI(vector_add, blocks, threads, 0, 0, Ad, Bd, Cd, N); // or
vector_add<<<blocks, threads, 0, 0>>>(Ad, Bd, Cd, N);

hipMemcpy(Ch, Cd, sizeof(float) * N, hipMemcpyDeviceToHost);

hipFree(Ad); hipFree(Bd); hipFree(Cd);
free(Ah); free(Bh); free(Ch); free(Cref);

```

➤ Step 2: launch kernel



Kernel launching

```

const int N = 10000;
float *Ah, *Bh, *Ch, *Cref;
float *Ad, *Bd, *Cd;

Ah = (float *)malloc(N * sizeof(float)); Bh = (float *)malloc(N * sizeof(float));
Ch = (float *)malloc(N * sizeof(float)); Cref = (float *)malloc(N * sizeof(float));

for (int i = 0; i < N; i++) {
    Ah[i] = sin(i) * 2.3;
    Bh[i] = cos(i) * 1.1;
    Cref[i] = Ah[i] + Bh[i];
}

hipMalloc((void **)&Ad, N * sizeof(float));
hipMalloc((void **)&Bd, N * sizeof(float));
hipMalloc((void **)&Cd, N * sizeof(float));

hipMemcpy(Ad, Ah, sizeof(float) * N, hipMemcpyHostToDevice);
hipMemcpy(Bd, Bh, sizeof(float) * N, hipMemcpyHostToDevice);

dim3 blocks, threads;
threads = dim3(256, 1, 1);
blocks = dim3((N + 256 - 1) / 256, 1, 1);


// hipLaunchKernelGGI(vector_add, blocks, threads, 0, 0, Ad, Bd, Cd, N); // or
vector_add<<<blocks, threads, 0>>>(Ad, Bd, Cd, N);

hipMemcpy(Ch, Cd, sizeof(float) * N, hipMemcpyDeviceToHost);

hipFree(Ad); hipFree(Bd); hipFree(Cd);
free(Ah); free(Bh); free(Ch); free(Cref);

```

➤ Step 2: launch kernel



Kernel launching

```

const int N = 10000;
float *Ah, *Bh, *Ch, *Cref;
float *Ad, *Bd, *Cd;

Ah = (float *)malloc(N * sizeof(float)); Bh = (float *)malloc(N * sizeof(float));
Ch = (float *)malloc(N * sizeof(float)); Cref = (float *)malloc(N * sizeof(float));

for (int i = 0; i < N; i++) {
    Ah[i] = sin(i) * 2.3;
    Bh[i] = cos(i) * 1.1;
    Cref[i] = Ah[i] + Bh[i];
}

hipMalloc((void **)&Ad, N * sizeof(float));
hipMalloc((void **)&Bd, N * sizeof(float));
hipMalloc((void **)&Cd, N * sizeof(float));

hipMemcpy(Ad, Ah, sizeof(float) * N, hipMemcpyHostToDevice);
hipMemcpy(Bd, Bh, sizeof(float) * N, hipMemcpyHostToDevice);

dim3 blocks, threads;
threads = dim3(256, 1, 1);
blocks = dim3((N + 256 - 1) / 256, 1, 1);


// hipLaunchKernelGGL(vector_add, blocks, threads, 0, 0, Ad, Bd, Cd, N); // or
vector_add<<<blocks, threads, 0, 0>>>(Ad, Bd, Cd, N);

hipMemcpy(Ch, Cd, sizeof(float) * N, hipMemcpyDeviceToHost);

hipFree(Ad); hipFree(Bd); hipFree(Cd);
free(Ah); free(Bh); free(Ch); free(Cref);

```

➤ Step 2: launch kernel



✓ Acceptable launch:

kernel <<<1, 1>>>(...)

Kernel itself

```
__global__ void vector_add(float *A, float *B, float *C, int n) {  
    int tid = threadIdx.x + blockIdx.x * blockDim.x;  
    if (tid < n) {  
        C[tid] = A[tid] + B[tid];  
    }  
}
```

Kernel itself

```
__global__ void vector_add(float *A, float *B, float *C, int n) {  
    int tid = threadIdx.x + blockIdx.x * blockDim.x;  
    if (tid < n) {  
        C[tid] = A[tid] + B[tid];  
    }  
}
```

- `__global__` : run on D, called from H
- `__device__` : run on D, called from D
- `__host__` : run on H, called from H

Kernel itself

```
__global__ void vector_add(float *A, float *B, float *C, int n) {  
    int tid = threadIdx.x + blockIdx.x * blockDim.x;  
    if (tid < n) {  
        C[tid] = A[tid] + B[tid];  
    }  
}
```

- `__global__` : run on D, called from H
- `__device__` : run on D, called from D
- `__host__` : run on H, called from H
- Global: always **void**

Kernel itself

```
__global__ void vector_add(float *A, float *B, float *C, int n) {  
    int tid = threadIdx.x + blockIdx.x * blockDim.x;  
    if (tid < n) {  
        C[tid] = A[tid] + B[tid];  
    }  
}
```

- `__global__` : run on D, called from H
- `__device__` : run on D, called from D
- `__host__` : run on H, called from H
- Global: always **void**
- Find global thread index

Kernel itself

```
__global__ void vector_add(float *A, float *B, float *C, int n) {  
    int tid = threadIdx.x + blockIdx.x * blockDim.x;  
    if (tid < n) {  
        C[tid] = A[tid] + B[tid];  
    }  
}
```

- `__global__` : run on D, called from H
- `__device__` : run on D, called from D
- `__host__` : run on H, called from H
- Global: always **void**
- Find global thread index
- Out of range check

Kernel itself

```
__global__ void vector_add(float *A, float *B, float *C, int n) {  
    int tid = threadIdx.x + blockIdx.x * blockDim.x;  
    if (tid < n) {  
        C[tid] = A[tid] + B[tid];  
    }  
}
```

- `__global__` : run on D, called from H
- `__device__` : run on D, called from D
- `__host__` : run on H, called from H
- Global: always **void**
- Find global thread index
- Out of range check
- Each thread calculates one sum

Kernel itself

```
__global__ void vector_add(float *A, float *B, float *C, int n) {  
  
    int tid = threadIdx.x + blockIdx.x * blockDim.x;  
    if (tid < n) {  
        C[tid] = A[tid] + B[tid];  
    }  
}
```

- `__global__` : run on D, called from H
- `__device__` : run on D, called from D
- `__host__` : run on H, called from H
- Global: always **void**
- Find global thread index
- Out of range check
- Each thread calculates one sum

Alternative: number of threads < number of tasks:

Kernel itself

```
__global__ void vector_add(float *A, float *B, float *C, int n) {  
  
    int tid = threadIdx.x + blockIdx.x * blockDim.x;  
    if (tid < n) {  
        C[tid] = A[tid] + B[tid];  
    }  
}
```

- `__global__` : run on D, called from H
- `__device__` : run on D, called from D
- `__host__` : run on H, called from H


- Global: always **void**

- Find global thread index

- Out of range check

- Each thread calculates one sum

Alternative: number of threads < number of tasks:




Kernel itself

```
__global__ void vector_add(float *A, float *B, float *C, int n) {  
  
    int tid = threadIdx.x + blockIdx.x * blockDim.x;  
    if (tid < n) {  
        C[tid] = A[tid] + B[tid];  
    }  
}
```

- `__global__` : run on D, called from H
- `__device__` : run on D, called from D
- `__host__` : run on H, called from H
- Global: always **void**
- Find global thread index
- Out of range check
- Each thread calculates one sum

Alternative: number of threads < number of tasks:




- Loop: first iteration

Kernel itself

```
__global__ void vector_add(float *A, float *B, float *C, int n) {  
  
    int tid = threadIdx.x + blockIdx.x * blockDim.x;  
    if (tid < n) {  
        C[tid] = A[tid] + B[tid];  
    }  
}
```

- `__global__` : run on D, called from H
- `__device__` : run on D, called from D
- `__host__` : run on H, called from H
- Global: always **void**
- Find global thread index
- Out of range check
- Each thread calculates one sum

Alternative: number of threads < number of tasks:




Kernel itself

```
_global__ void vector_add(float *A, float *B, float *C, int n) {
    int tid = threadIdx.x + blockIdx.x * blockDim.x;
    if (tid < n) {
        C[tid] = A[tid] + B[tid];
    }
}
```

- `_global__` : run on D, called from H
- `_device__` : run on D, called from D
- `_host__` : run on H, called from H
- Global: always **void**
- Find global thread index
- Out of range check
- Each thread calculates one sum


Alternative: number of threads < number of tasks:



```
_global__ void vector_add(float *A, float *B, float *C, int n) {
    int tid = threadIdx.x + blockIdx.x * blockDim.x;
    while (tid < n) {
        C[tid] = A[tid] + B[tid];
        tid += gridDim.x*blockDim.x;
    }
}
```

- **Thread linear addressing**: grid-size stride between passes through the loop

Shared memory




Key concepts

Optimizations


Honorable mentions

Shared memory



✓ **Hardware**: provided by SM, split between blocks on it


Shared memory



✓ **Hardware**: provided by SM, split between blocks on it


✓ Faster

Shared memory




- ✓ **Hardware**: provided by SM, split between blocks on it
- ✓ Faster
- ✓ Accessible to all threads in a block

Shared memory




- ✓ **Hardware:** provided by SM, split between blocks on it
- ✓ Faster
- ✓ Accessible to all threads in a block
- ✓ Limited size: too much requested -> only one active thread block

Shared memory



- ✓ **Hardware**: provided by SM, split between blocks on it
- ✓ Faster
- ✓ Accessible to all threads in a block
- ✓ Limited size: too much requested -> only one active thread block
- ✓ Does not require **coalescing**

Shared memory




➤ Static

```
__shared__ float data[256];
```

Declared inside kernel;
Size known at compile time

- ✓ **Hardware**: provided by SM, split between blocks on it
- ✓ Faster
- ✓ Accessible to all threads in a block
- ✓ Limited size: too much requested -> only one active thread block
- ✓ Does not require **coalescing**

Shared memory



➤ Static

```
__shared__ float data[256];
```

Declared inside kernel;
Size known at compile time

- ✓ **Hardware**: provided by SM, split between blocks on it
- ✓ Faster
- ✓ Accessible to all threads in a block
- ✓ Limited size: too much requested -> only one active thread block
- ✓ Does not require **coalescing**

➤ Dynamic

```
extern __shared__ float data[];
```


Size passed through kernel launch parameter **smem** ;
More flexible, but
Having more than one requires care

Parallel reduction

- Reductions: **aggregating elements** of an array **into a single value** through operations like summing, finding maximum or minimum, or performing logical operations


Parallel reduction

- Reductions: **aggregating elements** of an array **into a single value** through operations like summing, finding maximum or minimum, or performing logical operations
- Example: sum all elements of vector A



Parallel reduction

- Reductions: **aggregating elements** of an array **into a single value** through operations like summing, finding maximum or minimum, or performing logical operations
- Example: sum all N elements of vector A




Solution 1: loop through elements

Problem: slow

Parallel reduction

- Reductions: **aggregating elements** of an array **into a single value** through operations like summing, finding maximum or minimum, or performing logical operations
- Example: sum all N elements of vector A



Solution 2: loop through kernel launches halving number of threads

```
for(int m = N/2; m>0; m /= 2) {
    int threads = std::min(256,m);
    int blocks = std::max(m/256,1);
    reduce<<<blocks, threads>>>(A, m);
}
```

where:


```
__global__ void reduce(float *a, int m) {
    int tid = blockDim.x*blockIdx.x+threadIdx.x;
    a[tid] += a[tid+m];
}
```

3 calls to global memory,
doubled because of
iterating

Problem: memory bound, N must be power of 2

Parallel reduction

- Reductions: **aggregating elements** of an array **into a single value** through operations like summing, finding maximum or minimum, or performing logical operations
- Example: sum all N elements of vector A



Solution 3: iterate inside the kernel, accumulate partial sums there; use grid stride inside kernel

```
__global__ void reduce (float *a, int N) {
    int tid = blockDim.x*blockIdx.x+threadIdx.x;
    float tsum = 0.0f;
    for(int k=tid; k<N; k += gridDim.x*blockDim.x) tsum += a[k];
    a[tid] = tsum;
}
...
reduce<<< blocks, threads >>>(A,N);
reduce<<< 1, threads >>>(A,blocks*threads);
reduce<<< 1,1 >>>(A, threads);
```

Can we do better? Let's make threads cooperate

Parallel reduction

- Blocks will sum their individual accumulated totals using shared memory

```
__global__ void reduce(float * out, float * a, int N) {  
extern __shared__ float tsum[];  
int id = threadIdx.x; ← To save all partial sums  
in the corresponding  
part of shared memory  
int tid = blockDim.x*blockIdx.x+threadIdx.x;  
int stride = gridDim.x*blockDim.x;  
tsum[id] = 0.0;  
for(int k=tid;k<N;k+=stride) tsum[id] += a[k];  
__syncthreads();  
  
if(id<256 && id+256 < blockDim.x) tsum[id] += tsum[id+256]; __syncthreads();  
if(id<128) tsum[id] += tsum[id+128]; __syncthreads();  
if(id< 64) tsum[id] += tsum[id+ 64]; __syncthreads();  
if(id< 32) tsum[id] += tsum[id+ 32]; __syncthreads();  
  
if(id< 16) tsum[id] += tsum[id+16]; __syncwarp();  
if(id< 8) tsum[id] += tsum[id+ 8]; __syncwarp();  
if(id< 4) tsum[id] += tsum[id+ 4]; __syncwarp();  
if(id< 2) tsum[id] += tsum[id+ 2]; __syncwarp();  
if(id==0) out[blockIdx.x] = tsum[0]+tsum[1];  
}
```

Parallel reduction

- Blocks will sum their individual accumulated totals using shared memory

```
__global__ void reduce(float * out, float * a, int N) {  
extern __shared__ float tsum[];  
int id = threadIdx.x;  
int tid = blockDim.x*blockIdx.x+threadIdx.x;  
int stride = gridDim.x*blockDim.x;  
tsum[id] = 0.0;  
for(int k=tid;k<N;k+=stride) tsum[id] += a[k];  
__syncthreads();  
  
if(id<256 && id+256 < blockDim.x) tsum[id] += tsum[id+256]; __syncthreads();  
if(id<128) tsum[id] += tsum[id+128]; __syncthreads();  
if(id< 64) tsum[id] += tsum[id+ 64]; __syncthreads();  
if(id< 32) tsum[id] += tsum[id+ 32]; __syncthreads();  
  
if(id< 16) tsum[id] += tsum[id+16]; __syncwarp();  
if(id< 8) tsum[id] += tsum[id+ 8]; __syncwarp();  
if(id< 4) tsum[id] += tsum[id+ 4]; __syncwarp();  
if(id< 2) tsum[id] += tsum[id+ 2]; __syncwarp();  
if(id==0) out[blockIdx.x] = tsum[0]+tsum[1];  
}
```

Actually save all partial
sums in the
corresponding part of
shared memory

Parallel reduction

- Blocks will sum their individual accumulated totals using shared memory

```
__global__ void reduce(float * out, float * a, int N) {  
extern __shared__ float tsum[];  
int id = threadIdx.x;  
int tid = blockDim.x*blockIdx.x+threadIdx.x;  
int stride = gridDim.x*blockDim.x;  
tsum[id] = 0.0;  
for(int k=tid;k<N;k+=stride) tsum[id] += a[k];  
__syncthreads();  
  
if(id<256 && id+256 < blockDim.x) tsum[id] += tsum[id+256]; __syncthreads();  
if(id<128) tsum[id] += tsum[id+128]; __syncthreads();  
if(id< 64) tsum[id] += tsum[id+ 64]; __syncthreads();  
if(id< 32) tsum[id] += tsum[id+ 32]; __syncthreads();  
  
if(id< 16) tsum[id] += tsum[id+16]; __syncwarp();  
if(id< 8) tsum[id] += tsum[id+ 8]; __syncwarp();  
if(id< 4) tsum[id] += tsum[id+ 4]; __syncwarp();  
if(id< 2) tsum[id] += tsum[id+ 2]; __syncwarp();  
if(id==0) out[blockIdx.x] = tsum[0]+tsum[1];  
}
```

Do the reduction INSIDE shared memory (threads = 512 assumed)

Parallel reduction

- Blocks will sum their individual accumulated totals using shared memory

```
__global__ void reduce(float * out, float * a, int N) {  
extern __shared__ float tsum[];  
int id = threadIdx.x;  
int tid = blockDim.x*blockIdx.x+threadIdx.x;  
int stride = gridDim.x*blockDim.x;  
tsum[id] = 0.0;  
for(int k=tid;k<N;k+=stride) tsum[id] += a[k];  
__syncthreads();  
  
if(id<256 && id+256 < blockDim.x) tsum[id] += tsum[id+256]; __syncthreads();  
if(id<128) tsum[id] += tsum[id+128]; __syncthreads();  
if(id< 64) tsum[id] += tsum[id+ 64]; __syncthreads();  
if(id< 32) tsum[id] += tsum[id+ 32]; __syncthreads();  
  
if(id< 16) tsum[id] += tsum[id+16]; __syncwarp(); }  
if(id< 8) tsum[id] += tsum[id+ 8]; __syncwarp(); }  
if(id< 4) tsum[id] += tsum[id+ 4]; __syncwarp(); }  
if(id< 2) tsum[id] += tsum[id+ 2]; __syncwarp(); }  
if(id==0) out[blockIdx.x] = tsum[0]+tsum[1]; }
```

Still reduction in shared memory, but making use of warps (for AMD we would start `_syncwarp` the line above)

Parallel reduction

- Blocks will sum their individual accumulated totals using shared memory

```
__global__ void reduce(float * out, float * a, int N) {  
extern __shared__ float tsum[];  
int id = threadIdx.x;  
int tid = blockDim.x*blockIdx.x+threadIdx.x;  
int stride = gridDim.x*blockDim.x;  
tsum[id] = 0.0;  
for(int k=tid;k<N;k+=stride) tsum[id] += a[k];  
__syncthreads();  
  
if(id<256 && id+256 < blockDim.x) tsum[id] += tsum[id+256]; __syncthreads();  
if(id<128) tsum[id] += tsum[id+128]; __syncthreads();  
if(id< 64) tsum[id] += tsum[id+ 64]; __syncthreads();  
if(id< 32) tsum[id] += tsum[id+ 32]; __syncthreads();  
  
if(id< 16) tsum[id] += tsum[id+16]; __syncwarp();  
if(id< 8) tsum[id] += tsum[id+ 8]; __syncwarp();  
if(id< 4) tsum[id] += tsum[id+ 4]; __syncwarp();  
if(id< 2) tsum[id] += tsum[id+ 2]; __syncwarp();  
if(id==0) out[blockIdx.x] = tsum[0]+tsum[1];  
}
```


Further improvements:

- ✓ Make threads a template
- ✓ Use cooperative groups
- ✓ Use warp reduction
- ✓ Halve the number of blocks and "double load"


Some good sources:

<https://developer.download.nvidia.com/assets/cuda/files/reduction.pdf>
https://github.com/NVIDIA/cuda-samples/blob/master/Samples/2_Concepts_and_Techniques/reduction/
<https://developer.nvidia.com/blog/faster-parallel-reductions-kepler>

Matrix transpose





Matrix transpose



➤ CPU:

```
//Bx = Ay; By = Ax;  
for(int iy=0; iy<Ay; iy++){  
    for(int ix=0; ix<Ax; ix++){  
        B[ix*Ay+iy] = A[iy*Ax+ix];  
    }  
}
```

Matrix transpose




- CPU:


```
//Bx = Ay; By = Ax;
for(int iy=0; iy<Ay; iy++){
    for(int ix=0; ix<Ax; ix++){
        B[ix*Ay+iy] = A[iy*Ax+ix];
    }
}
```
- Naïve version on GPU:


```
__global__ void gpu_transpose(const float * A, float * B, int Ax, int Ay){
int ix = blockIdx.x*blockDim.x + threadIdx.x;
int iy = blockIdx.y*blockDim.y + threadIdx.y;
if(iy >= Ay || ix >= Ax) return;
    B[ix * Ay + iy] = A[iy * Ax + ix];
}
...
dim3 threads ={threadIdx.x, threadIdx.y, 1};
dim3 blocks ={(Ax+threads.x-1)/threads.x, (Ay+threads.y-1)/threads.y, 1};
gpu_transpose<<<blocks,threads>>>(A, B, Ax, Ay);

//Ax = 5; Ay = 8;
//threadIdx.x = 4; threadIdx.y = 2;
//blockIdx.x = 2; blockIdx.y = 4;
```

Matrix transpose



➤ CPU:

```
//Bx = Ay; By = Ax;
for(int iy=0; iy<Ay; iy++){
    for(int ix=0; ix<Ax; ix++){
        B[ix*Ay+iy] = A[iy*Ax+ix];
    }
}
```

➤ Naïve version on GPU:


```
__global__ void gpu_transpose(const float * A, float * B, int Ax, int Ay){
int ix = blockIdx.x*blockDim.x + threadIdx.x;
int iy = blockIdx.y*blockDim.y + threadIdx.y;
if(iy >= Ay || ix >= Ax) return;

B[ix * Ay + iy] = A[iy * Ax + ix];
}

...
dim3 threads ={threadIdx.x, threadIdx.y, 1};
dim3 blocks ={(Ax+threads.x-1)/threads.x, (Ay+threads.y-1)/threads.y, 1};
gpu_transpose<<<blocks,threads>>>(A, B, Ax, Ay);

//Ax = 5; Ay = 8;
//threadIdx.x = 4; threadIdx.y = 2;
//blockIdx.x = 2; blockIdx.y = 4;
```

Matrix transpose



- CPU:

```
//Bx = Ay; By = Ax;
for(int iy=0; iy<Ay; iy++){
    for(int ix=0; ix<Ax; ix++){
        B[ix*Ay+iy] = A[iy*Ax+ix];
    }
}
```

- Naïve version on GPU:


```
__global__ void gpu_transpose(const float * A, float * B, int Ax, int Ay){
int ix = blockIdx.x*blockDim.x + threadIdx.x;
int iy = blockIdx.y*blockDim.y + threadIdx.y;
if(iy >= Ay || ix >= Ax) return;

B[ix * Ay + iy] = A[iy * Ax + ix];
}

...
dim3 threads ={threadIdx.x, threadIdx.y, 1};
dim3 blocks ={(Ax+threads.x-1)/threads.x, (Ay+threads.y-1)/threads.y, 1};
gpu_transpose<<<blocks,threads>>>(A, B, Ax, Ay);

//Ax = 5; Ay = 8;
//threadIdx.x = 4; threadIdx.y = 2;
//blockIdx.x = 2; blockIdx.y = 4;
```

Matrix transpose



➤ Naïve version on GPU:

```
__global__ void gpu_transpose(const float * A, float * B, int Ax, int Ay){
int ix = blockIdx.x*blockDim.x + threadIdx.x;
int iy = blockIdx.y*blockDim.y + threadIdx.y;
if(iy >= Ay || ix >= Ax) return;

B[ix * Ay + iy] = A[iy * Ax + ix];
}


...
dim3 threads ={threadIdx, thready, 1};
dim3 blocks ={(Ax+threads.x-1)/threads.x, (Ay+threads.y-1)/threads.y, 1};
gpu_transpose<<<blocks,threads>>>(A, B, Ax, Ay);

//Ax = 5; Ay = 8;
//threadx = 4; thready = 2;
//blockx = 2; blocky = 4;
```

➤ CPU:

```
//Bx = Ay; By = Ax;  
for(int iy=0; iy<Ay; iy++){  
    for(int ix=0; ix<Ax; ix++){  
        B[ix*Ay+iy] = A[iy*Ax+ix];  
    }  
}
```

Matrix transpose



This is coalescent:
nearest in memory
threads address
neighbouring elements
of A

- CPU:

```
//Bx = Ay; By = Ax;
for(int iy=0; iy<Ay; iy++){
    for(int ix=0; ix<Ax; ix++){
        B[ix*Ay+iy] = A[iy*Ax+ix];
    }
}
```
 - Naïve version on GPU:


```
__global__ void gpu_transpose(const float * A, float * B, int Ax, int Ay){
    int ix = blockIdx.x*blockDim.x + threadIdx.x;
    int iy = blockIdx.y*blockDim.y + threadIdx.y;
    if(iy >= Ay || ix >= Ax) return;

    B[ix * Ay + iy] = A[iy * Ax + ix];
}

...
dim3 threads ={threadIdx.x, threadIdx.y, 1};
dim3 blocks ={(Ax+threads.x-1)/threads.x, (Ay+threads.y-1)/threads.y, 1};
gpu_transpose<<<blocks,threads>>>(A, B, Ax, Ay);

//Ax = 5; Ay = 8;
//threadIdx.x = 4; threadIdx.y = 2;
//blockIdx.x = 2; blockIdx.y = 4;
```

Matrix transpose



➤ CPU:

```
//Bx = Ay; By = Ax;
for(int iy=0; iy<Ay; iy++){
    for(int ix=0; ix<Ax; ix++){
        B[ix*Ay+iy] = A[iy*Ax+ix];
    }
}
```

➤ Naïve version on GPU:

```
__global__ void gpu_transpose(const float * A, float * B, int Ax, int Ay){
int ix = blockIdx.x*blockDim.x + threadIdx.x;
int iy = blockIdx.y*blockDim.y + threadIdx.y;
if(iy >= Ay || ix >= Ax) return;

B[ix * Ay + iy] = A[iy * Ax + ix];
}

...
```


dim3 threads ={threadIdx.x, threadIdx.y, 1};
dim3 blocks ={(Ax+threads.x-1)/threads.x, (Ay+threads.y-1)/threads.y, 1};
gpu_transpose<<<blocks,threads>>>(A, B, Ax, Ay);

//Ax = 5; Ay = 8;
//threadIdx.x = 4; threadIdx.y = 2;
//blockIdx.x = 2; blockIdx.y = 4;

Matrix transpose

NxM

MxN



- Improved version on GPU – using shared memory:

```

template<int tile_dim>
__global__ void gpu_transpose_SM(const float* A, float* B, int Ax, int Ay){
    __shared__ float tile[tile_dim][tile_dim];

    int x = blockIdx.x * tile_dim + threadIdx.x; // column in A
    int y = blockIdx.y * tile_dim + threadIdx.y; // row in A
    int tx = threadIdx.x; // column in tile
    int ty = threadIdx.y; // row in tile

    if (x < Ax && y < Ay){ tile[ty][tx] = A[y * Ax + x]; }
    __syncthreads();

    int bx = blockIdx.y * TILE_DIM + threadIdx.x; //column in B
    int by = blockIdx.x * TILE_DIM + threadIdx.y; //row in B


    if (bx < Ay && by < Ax){ B[by * Ay + bx] = tile[tx][ty]; }

}

...
dim3 thread = {tile_dim, tile_dim, 1};
dim3 blocks = {(Ax+threads.x-1)/threads.x, (Ay+threads.y-1)/threads.y, 1};
gpu transpose S<tile dim><<<blocks, threads>>>(A, B, Ax, Ay);

```

Matrix transpose



- Improved version on GPU – using shared memory:

```
template<int tile_dim>
__global__ void gpu_transpose_SM(const float* A, float* B, int Ax, int Ay){
    __shared__ float tile[tile_dim][tile_dim];

    int x = blockIdx.x * tile_dim + threadIdx.x; // column in A
    int y = blockIdx.y * tile_dim + threadIdx.y; // row in A
    int tx = threadIdx.x; // column in tile
    int ty = threadIdx.y; // row in tile


    if (x < Ax && y < Ay){ tile[ty][tx] = A[y * Ax + x]; }
    __syncthreads();

    int bx = blockIdx.y * TILE_DIM + threadIdx.x; //column in B
    int by = blockIdx.x * TILE_DIM + threadIdx.y; //row in B

    if (bx < Ay && by < Ax){ B[by * Ay + bx] = tile[tx][ty]; }
}

...
dim3 thread = {tile_dim, tile_dim, 1};
dim3 blocks = {((Ax+threads.x-1)/threads.x, (Ay+threads.y-1)/threads.y, 1)};
gpu_transpose_S<tile_dim><<<blocks, threads>>>(A, B, Ax, Ay);
```

Matrix transpose



- Improved version on GPU – using shared memory:

```
template<int tile_dim>
__global__ void gpu_transpose_SM(const float* A, float* B, int Ax, int Ay){
    __shared__ float tile[tile_dim][tile_dim];

    int x = blockIdx.x * tile_dim + threadIdx.x; // column in A
    int y = blockIdx.y * tile_dim + threadIdx.y; // row in A
    int tx = threadIdx.x; // column in tile
    int ty = threadIdx.y; // row in tile

    if (x < Ax && y < Ay){ tile[ty][tx] = A[y * Ax + x]; }
    __syncthreads();

    int bx = blockIdx.y * TILE_DIM + threadIdx.x; //column in B
    int by = blockIdx.x * TILE_DIM + threadIdx.y; //row in B


    if (bx < Ay && by < Ax){ B[by * Ay + bx] = tile[tx][ty]; }
}

...
dim3 thread = {tile_dim, tile_dim, 1};
dim3 blocks = {((Ax+threads.x-1)/threads.x, (Ay+threads.y-1)/threads.y, 1};
gpu_transpose_S<tile_dim><<<blocks, threads>>>(A, B, Ax, Ay);
```


Matrix transpose

NxM

MxN



blockIdx.y = 0
blockIdx.x = 0



- Improved version on GPU – using shared memory:

```
template<int tile_dim>
__global__ void gpu_transpose_SM(const float* A, float* B, int Ax, int Ay){
    __shared__ float tile[tile_dim][tile_dim];

    int x = blockIdx.x * tile_dim + threadIdx.x; // column in A
    int y = blockIdx.y * tile_dim + threadIdx.y; // row in A
    int tx = threadIdx.x; // column in tile
    int ty = threadIdx.y; // row in tile

    if (x < Ax && y < Ay){ tile[ty][tx] = A[y * Ax + x]; }
    __syncthreads();

    int bx = blockIdx.y * TILE_DIM + threadIdx.x; //column in B
    int by = blockIdx.x * TILE_DIM + threadIdx.y; //row in B


    if (bx < Ay && by < Ax){ B[by * Ay + bx] = tile[tx][ty]; }

}

...
dim3 thread = {tile_dim, tile_dim, 1};
dim3 blocks = {((Ax+threads.x-1)/threads.x, (Ay+threads.y-1)/threads.y, 1)};
gpu_transpose_S<tile_dim><<<blocks, threads>>>(A, B, Ax, Ay);
```

Matrix transpose

NxM MxN



- Improved version on GPU – using shared memory:

```
template<int tile_dim>
__global__ void gpu_transpose_SM(const float* A, float* B, int Ax, int Ay){
    __shared__ float tile[tile_dim][tile_dim];

    int x = blockIdx.x * tile_dim + threadIdx.x; // column in A
    int y = blockIdx.y * tile_dim + threadIdx.y; // row in A
    int tx = threadIdx.x; // column in tile
    int ty = threadIdx.y; // row in tile

    if (x < Ax && y < Ay){ tile[ty][tx] = A[y * Ax + x]; }
    __syncthreads();

    int bx = blockIdx.y * TILE_DIM + threadIdx.x; //column in B
    int by = blockIdx.x * TILE_DIM + threadIdx.y; //row in B


    if (bx < Ay && by < Ax){ B[by * Ay + bx] = tile[tx][ty]; }

}


dim3 thread = {tile_dim, tile_dim, 1};
dim3 blocks = {((Ax+threads.x-1)/threads.x, (Ay+threads.y-1)/threads.y, 1)};
gpu_transpose_S<tile_dim><<<blocks, threads>>>(A, B, Ax, Ay);
```

Matrix transpose

NxM



MxN



- Improved version on GPU – using shared memory:

```
template<int tile_dim>
__global__ void gpu_transpose_SM(const float* A, float* B, int Ax, int Ay){
    __shared__ float tile[tile_dim][tile_dim];

    int x = blockIdx.x * tile_dim + threadIdx.x; // column in A
    int y = blockIdx.y * tile_dim + threadIdx.y; // row in A
    int tx = threadIdx.x; // column in tile
    int ty = threadIdx.y; // row in tile

    if (x < Ax && y < Ay){ tile[ty][tx] = A[y * Ax + x]; }
    __syncthreads();

    int bx = blockIdx.y * TILE_DIM + threadIdx.x; //column in B
    int by = blockIdx.x * TILE_DIM + threadIdx.y; //row in B

    if (bx < Ay && by < Ax){ B[by * Ay + bx] = tile[tx][ty]; }


}

dim3 thread = {tile_dim, tile_dim, 1};
dim3 blocks = {((Ax+threads.x-1)/threads.x, (Ay+threads.y-1)/threads.y, 1)};
gpu_transpose_S<tile_dim><<<blocks, threads>>>(A, B, Ax, Ay);
```

Matrix transpose

NxM

MxN



- Improved version on GPU – using shared memory:

```

template<int tile_dim>
__global__ void gpu_transpose_SM(const float* A, float* B, int Ax, int Ay){
    __shared__ float tile[tile_dim][tile_dim];

    int x = blockIdx.x * tile_dim + threadIdx.x; // column in A
    int y = blockIdx.y * tile_dim + threadIdx.y; // row in A
    int tx = threadIdx.x; // column in tile
    int ty = threadIdx.y; // row in tile

    if (x < Ax && y < Ay){ tile[ty][tx] = A[y * Ax + x]; }
    __syncthreads();

    int bx = blockIdx.y * TILE_DIM + threadIdx.x; //column in B
    int by = blockIdx.x * TILE_DIM + threadIdx.y; //row in B


    if (bx < Ay && by < Ax){ B[by * Ay + bx] = tile[tx][ty]; }

}

dim3 thread = {tile_dim, tile_dim, 1};
dim3 blocks = {((Ax+threads.x-1)/threads.x, (Ay+threads.y-1)/threads.y, 1)};
gpu_transpose_S<tile_dim><<<blocks, threads>>>(A, B, Ax, Ay);

```

Streams



- Allow to split data into pieces that can be processed and copied **simultaneously**
- Set with “stream” parameter inside kernel call
- Have to be accompanied with *hipMemcpyAsync*
- Default stream is 0

Useful for:

- Large “frame by frame” data processing
- Simultaneously running memory-bound and computationally-bound kernels

Data transfer options

- Zero-copy memory: pinned memory on CPU can be modified and accessed by GPU directly; device pointer is still needed, referring to the same memory
 - Unified memory: same pointer is used for CPU and GPU, memory transfers are “hidden”
-
- ✓ Extra care with synchronization: `hipDeviceSynchronize` and implicit memory copying
 - ✓ UM can be slow on windows
 - ✓ Zero-copy becomes slow upon repeated GPU read

... and more

- Cooperative groups: convenient functionality to get bid, tid etc; allows grouping threads in tiles
- Thrust library: Adds std::vector to GPU; has some useful features like complex numbers
- Multi-GPU: cudaSetDevice, cudaGetDeviceCount etc.; Unified Virtual Addressing – treats memory of devices as a unit
- Warp-level programming: Parallel reduction as an example; speed-up
- Atomic operations: Queue threads' requests to modify memory simultaneously

Look out for

- Coalescent memory access
- Avoid thread divergence within warp – will make them run in sets
- Avoid deadlocks – thread synchronisations in *if* statements
- Don't allow threads to read/write in the same place in memory – race conditions
- Check if you are trying to access out of bound memory
- Use error checks:

```
#define ASSERT_CUDA_SUCCESS(cudaCall)
{
    cudaError_t error = cudaCall;
    if(error != cudaSuccess) {
        std::fprintf(
            stderr, "Error on line %i, file %s: %s\n", __LINE__, __FILE__, cudaGetString(error));
        std::exit(EXIT_FAILURE);
    }
}
```

To take home message

- Learn about the hardware
- Keep track of memory allocations and flow
- Track data/threads indexing correspondence and coalescence
- Don't trust ChatGPT too much
- If in doubt, opt for an easier solution and make GPU communicate back to you

<https://github.com/ENCCS/gpu-programming>